


Generative Adversarial Network



CS-EJ3311 - Deep Learning with Python
24.10.-11.12.2022
Aalto University & FiTech.io

5.12.2022 Shamsi Abdurakhmanova

Task A (source)

Domain $\mathcal{D}_s = \{\mathcal{X}_s, P(\mathcal{X}_s)\}$

Task $\mathcal{T}_s = \{\mathcal{Y}_s, P(\mathcal{Y}_s|\mathcal{X}_s)\}$

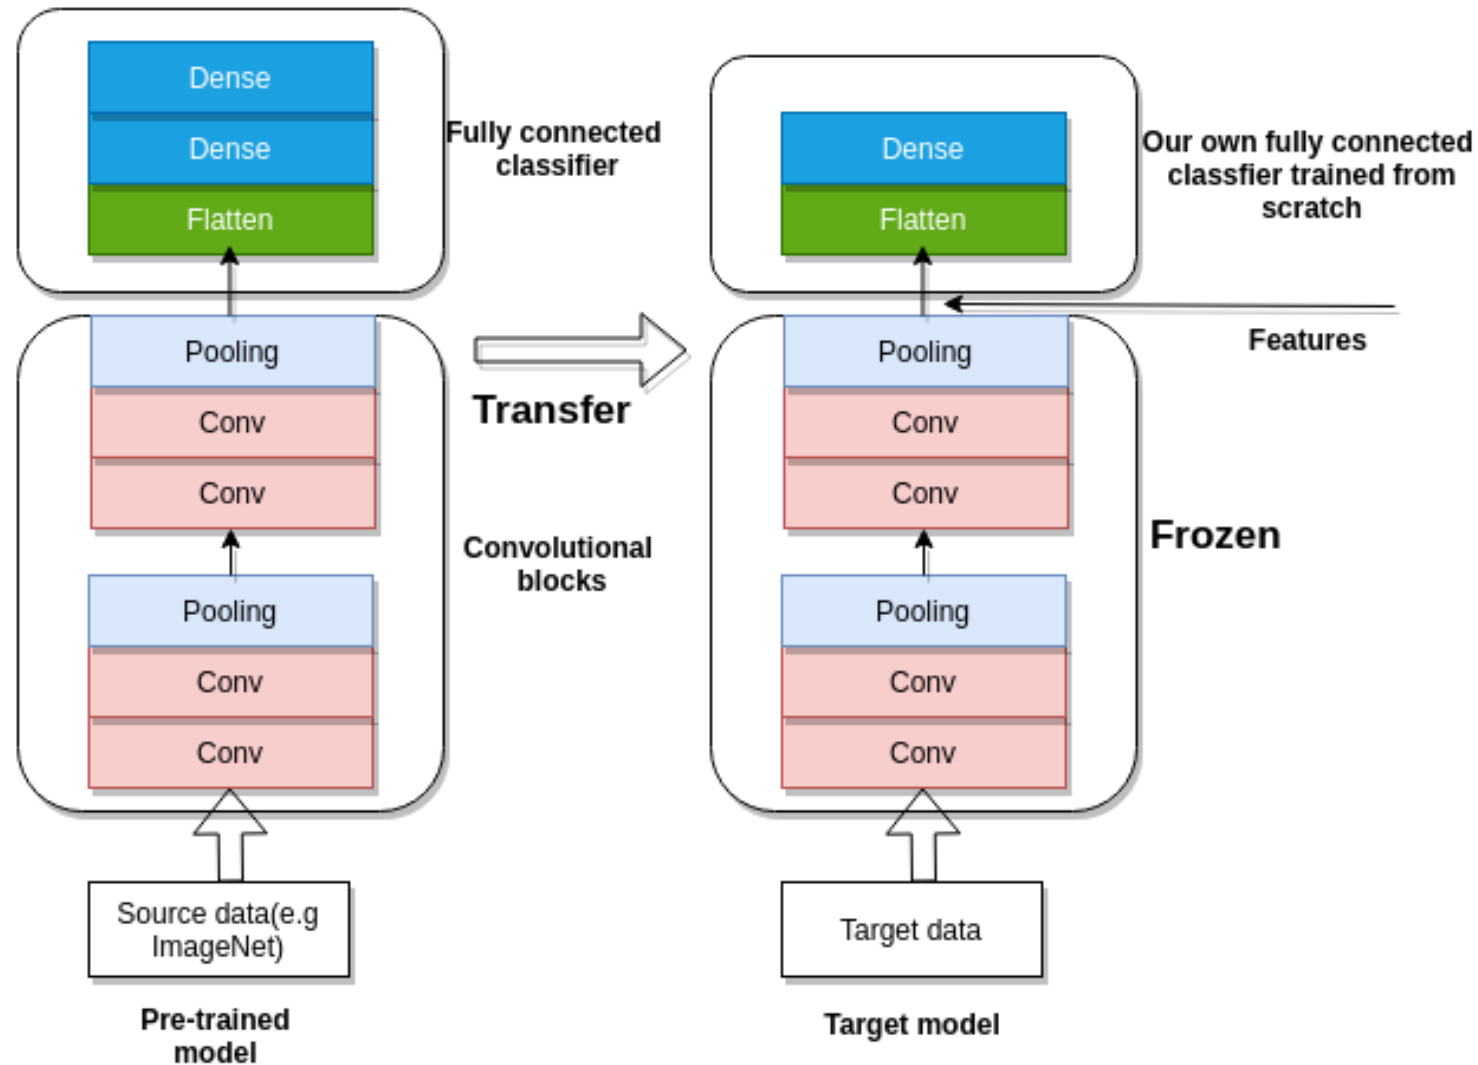
Task B (target)

Domain $\mathcal{D}_T = \{\mathcal{X}_T, P(\mathcal{X}_T)\}$

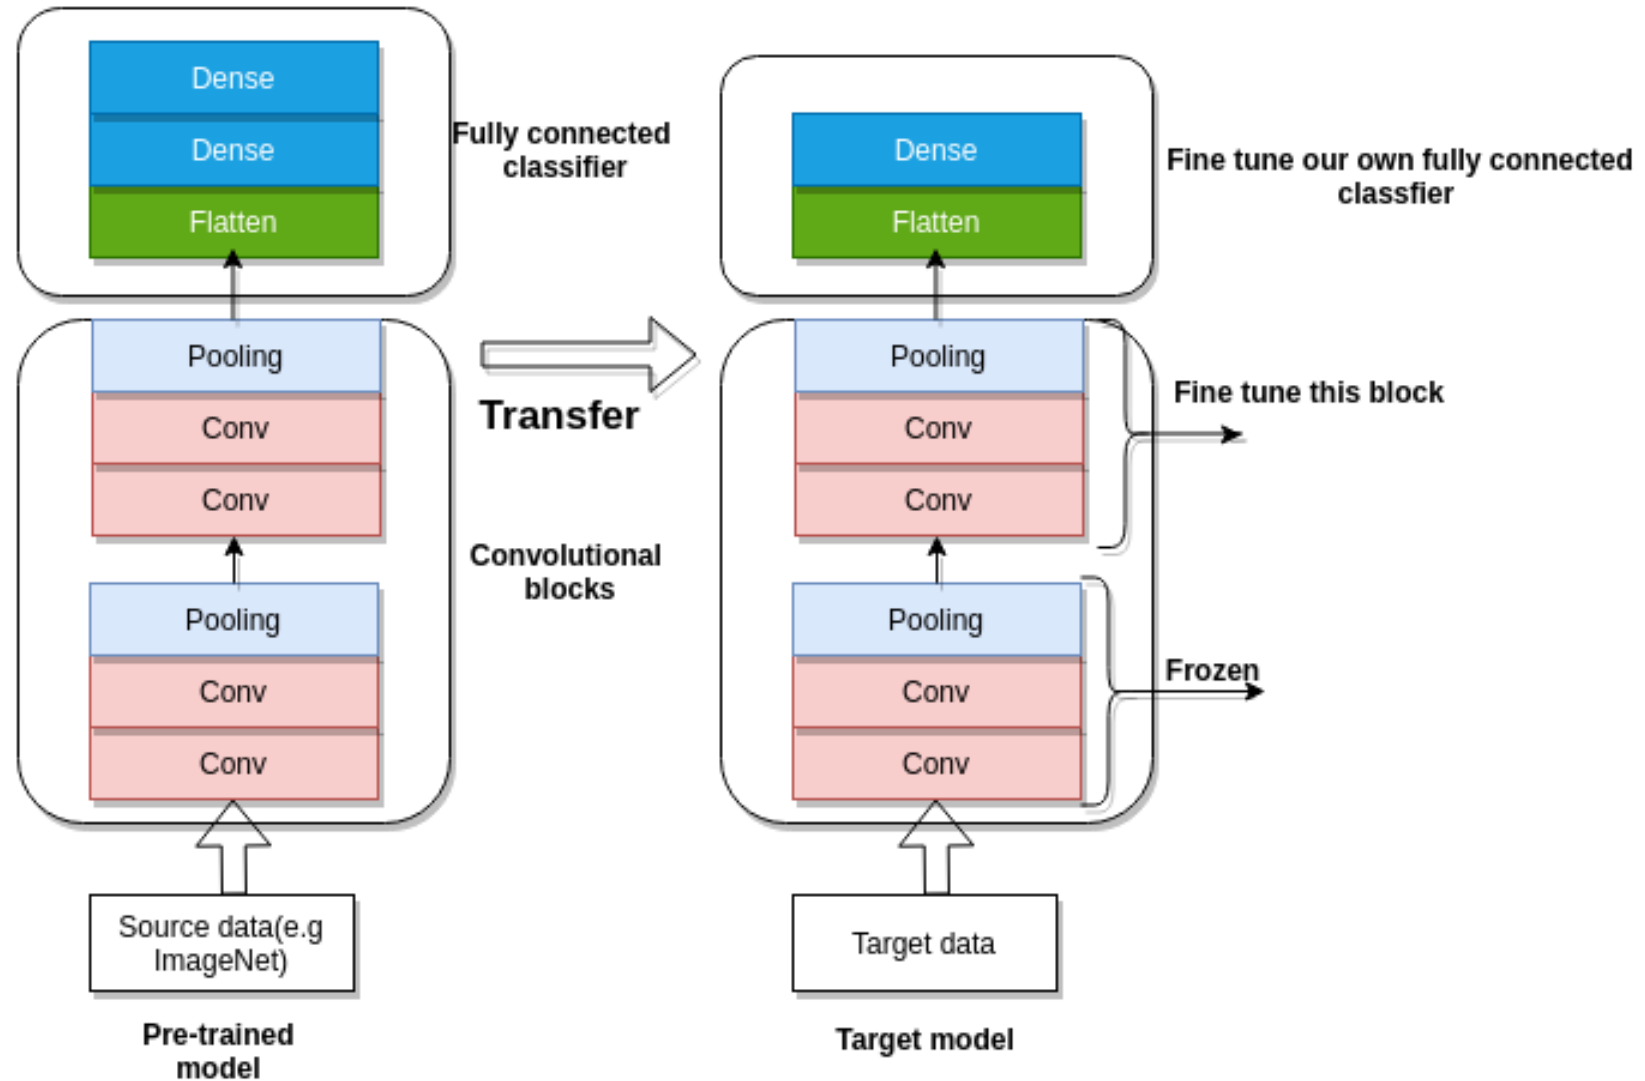
Task $\mathcal{T}_T = \{\mathcal{Y}_T, P(\mathcal{Y}_T|\mathcal{X}_T)\}$

Transfer learning – learn $P(\mathcal{Y}_T|\mathcal{X}_T)$ using knowledge $\mathcal{D}_s, \mathcal{T}_s$

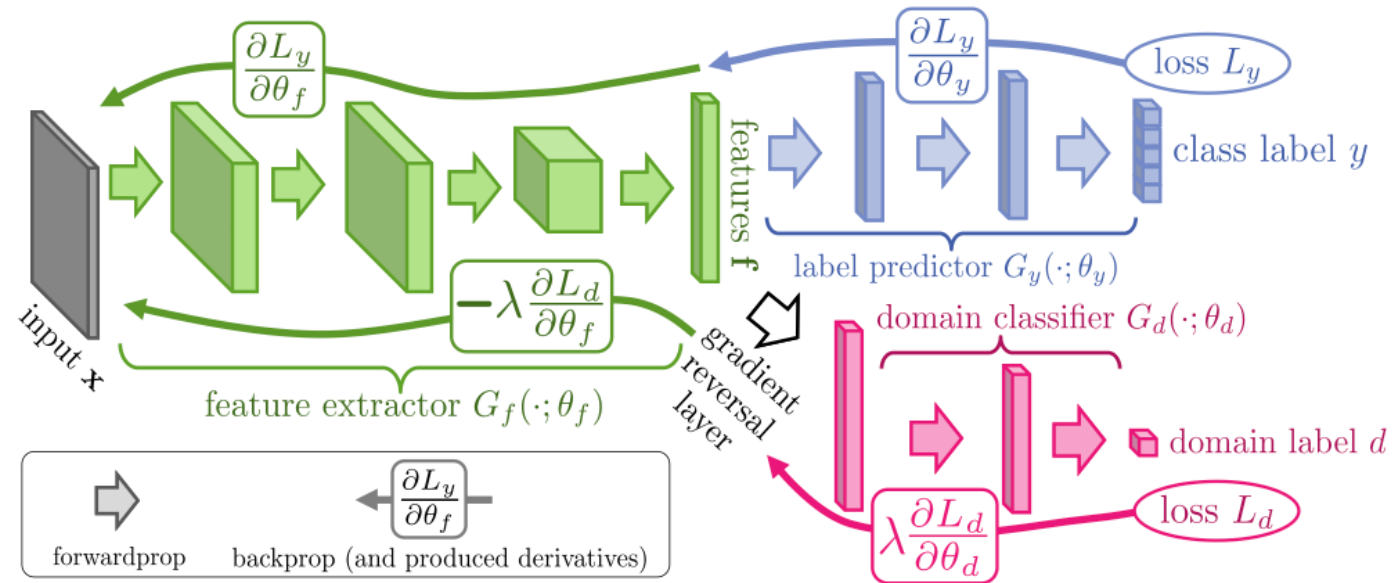


Pre-trained model for feature extraction



Fine-tuning pre-trained model





Transfer learning when no labels are available




0
1
2
...
9

MNIST
MNIST-M

Image classifying ANN



Data (image) generating ANN






generated sample

targer sample



LOSE



"Fake!"



"Real!"


WIN




Generator
"Artist"

WIN

generated sample



training sample




"Real!"




Discriminator
"Critic"

LOSE


"Real!"




a)



b)



c)



d)

“Generative Adversarial Nets” Ian Goodfellow 2014

GAN PROGRESS ON FACE GENERATION

Source: Goodfellow et al., 2014; Radford et al., 2016; Liu & Tuzel, 2016; Karras et al., 2018; Karras et al., 2019; Goodfellow, 2019; Karras et al., 2020; AI Index, 2021





Figure 2.1.7

<https://thispersondoesnotexist.com>

Input



Output



Input



Output




horse \rightarrow zebra

Input



Output



zebra \rightarrow horse





Random
vector

Generator

Discriminator

Fake or
Real






GAN

```
# random vector  
codings_size=...
```

```
# sample random vectors  
noise = tf.random.normal(shape=[batch_size, codings_size])
```




```
gen_out_size=... # shape of the generated sample
```



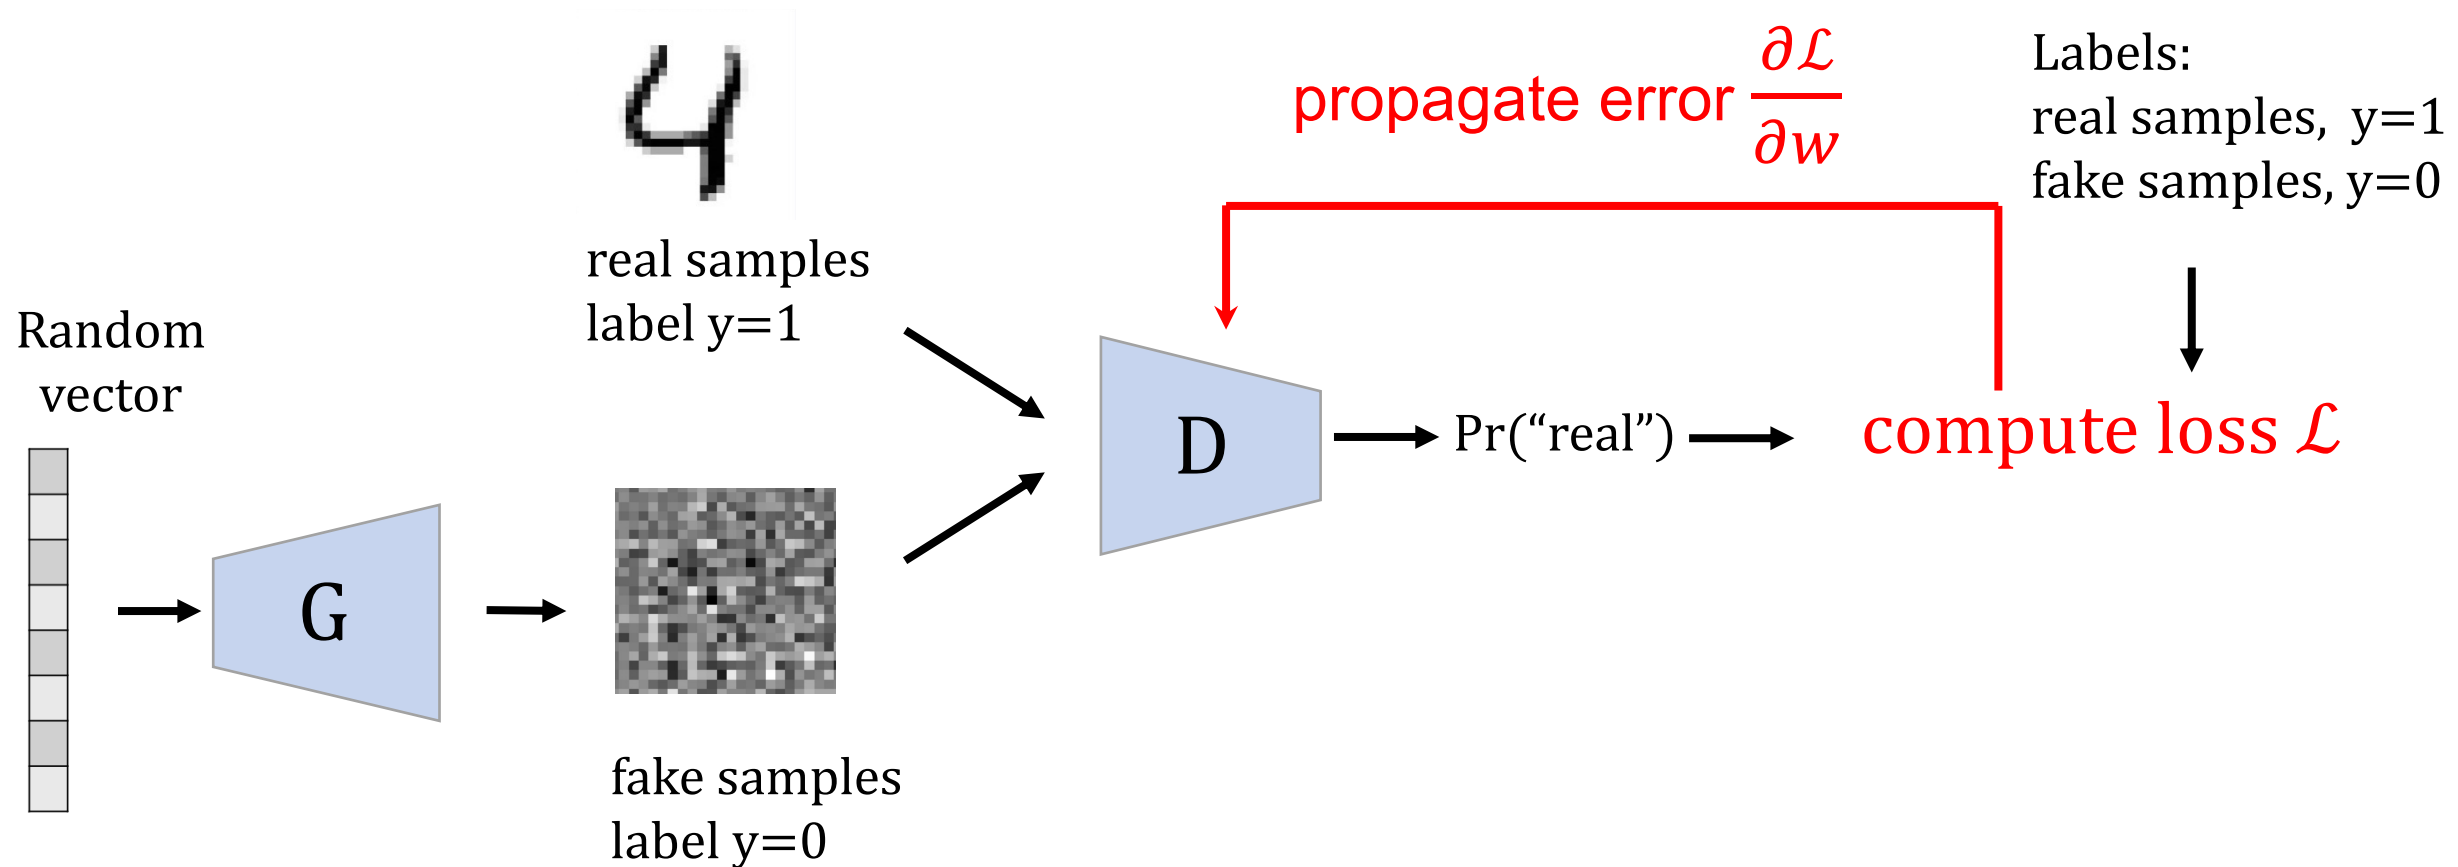

```
generator = tf.keras.models.Sequential([  
    tf.keras.layers.Dense(..., input_shape=[codings_size]),  
    tf.keras.layers.Dense(...),  
    tf.keras.layers.Dense(gen_out_size)  
])
```

```
discriminator = tf.keras.models.Sequential([  
    tf.keras.layers.Dense(..., input_shape=[gen_out_size ]),  
    tf.keras.layers.Dense(...),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```


```
gan = tf.keras.models.Sequential([generator, discriminator])
```



Phase 1. Discriminator training



Phase 2. Generator training



Discriminator

maximize loss:


$$\log D \left(\mathbf{x}^{(i)} \right) + \log \left(1 - D \left(G \left(\mathbf{z}^{(i)} \right) \right) \right)$$

minimize
binary CE

$$L = -\frac{1}{m} \sum_{i=1}^m (y_i \cdot \log(\hat{y}_i) + (1 - y_i) \cdot \log(1 - \hat{y}_i))$$

$x^{(i)}$ - real data point, assign label $y = 1$

$G(z^{(i)})$ - generated “fake” data point, assign label $y = 0$



Generator
minimize loss:

$$\log \left(1 - D \left(G \left(\mathbf{z}^{(i)} \right) \right) \right)$$

minimize
binary CE

$$L = -\frac{1}{m} \sum_{i=1}^m y_i \cdot \log(\hat{y}_i)$$

$G(\mathbf{z}^{(i)})$ – generated “fake” data point, assign lable **$y = 1$**

build one model

```
gan = tf.keras.models.Sequential([generator, discriminator])
```

compile discriminator

discriminator.trainable=True when training discriminator

```
discriminator.compile(loss="binary_crossentropy",  
optimizer="rmsprop")
```

compile gan

discriminator.trainable=False when training gan

```
discriminator.trainable = False
```

```
gan.compile(loss="binary_crossentropy", optimizer="rmsprop")
```

```
def train_gan(gan, dataset, batch_size, codings_size, n_epochs=90):
```

```
    generator, discriminator = gan.layers
```

```
    for epoch in range(n_epochs):
```

```
        for X_batch in dataset:
```

```
            # phase 1 - training the discriminator
```

random vector

```
            noise = tf.random.normal(shape=[batch_size, codings_size])
```

```
            gen_samples = generator(noise)
```

fake samples

```
            X_fake_and_real = tf.concat([gen_samples, tf.cast(X_batch, tf.float32)], axis=0)
```

```
            y1 = tf.constant([[0.]] * batch_size + [[1.]] * batch_size)
```

labels:

```
            discriminator.train_on_batch(X_fake_and_real, y1)
```

real y=1

```
            # phase 2 - training the generator
```

fake y=0

```
            noise = tf.random.normal(shape=[batch_size, codings_size])
```

random vector


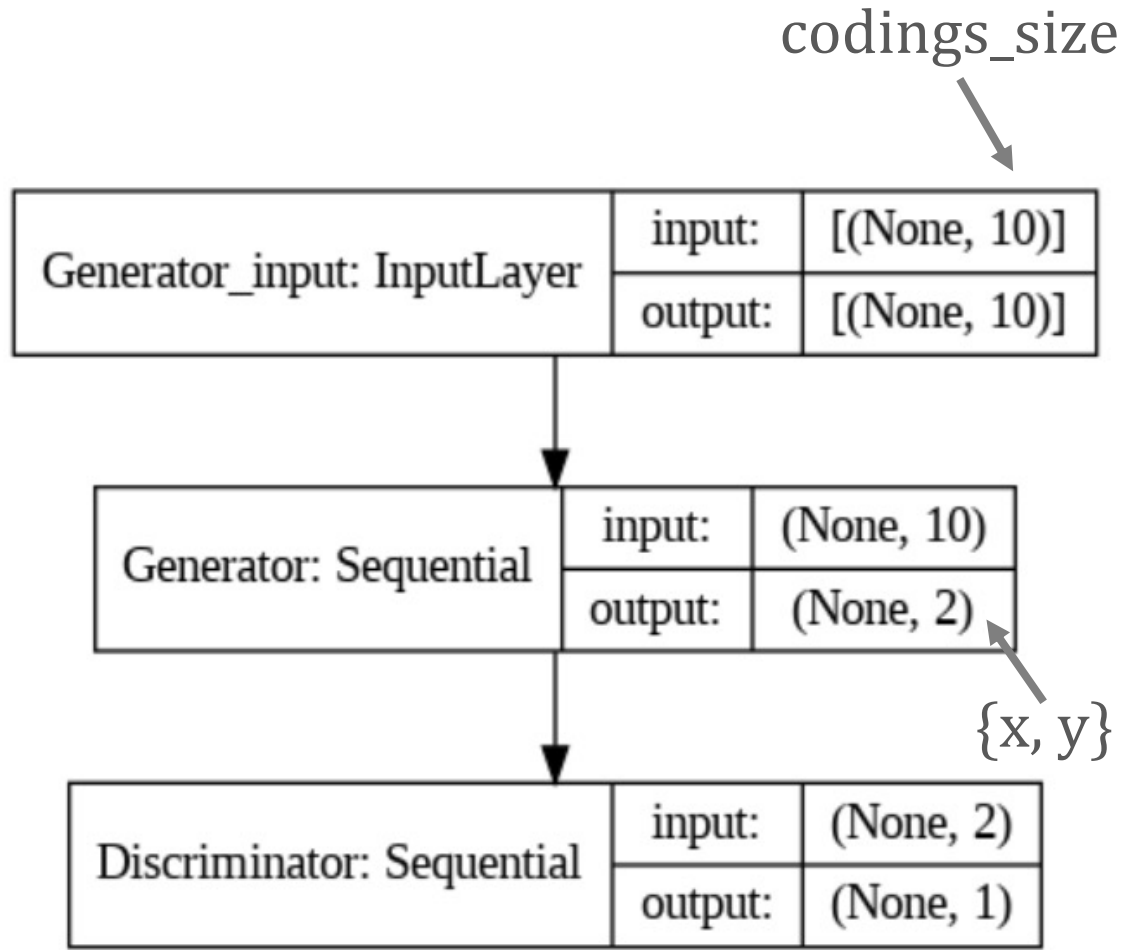
```
            y2 = tf.constant([[1.]] * batch_size)
```

labels:

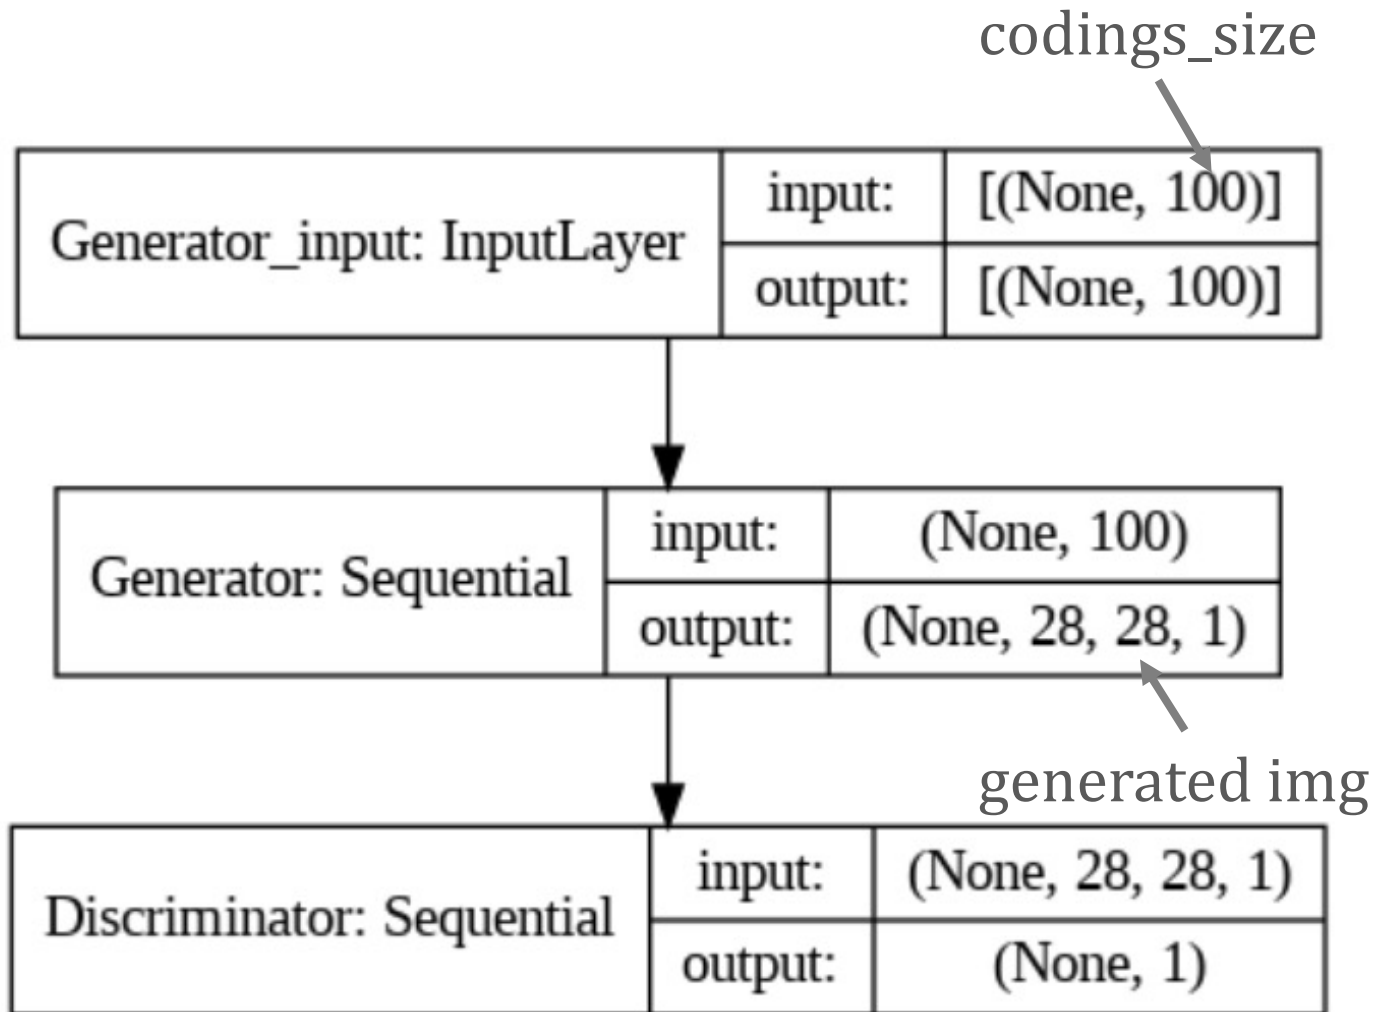
```
            gan.train_on_batch(noise, y2)
```

fake y=1


Generate samples from quadratic distribution.




DCGAN on MNIST dataset



Epoch 1/2




Epoch 2/2




DCGAN – Deep Convolutional GAN

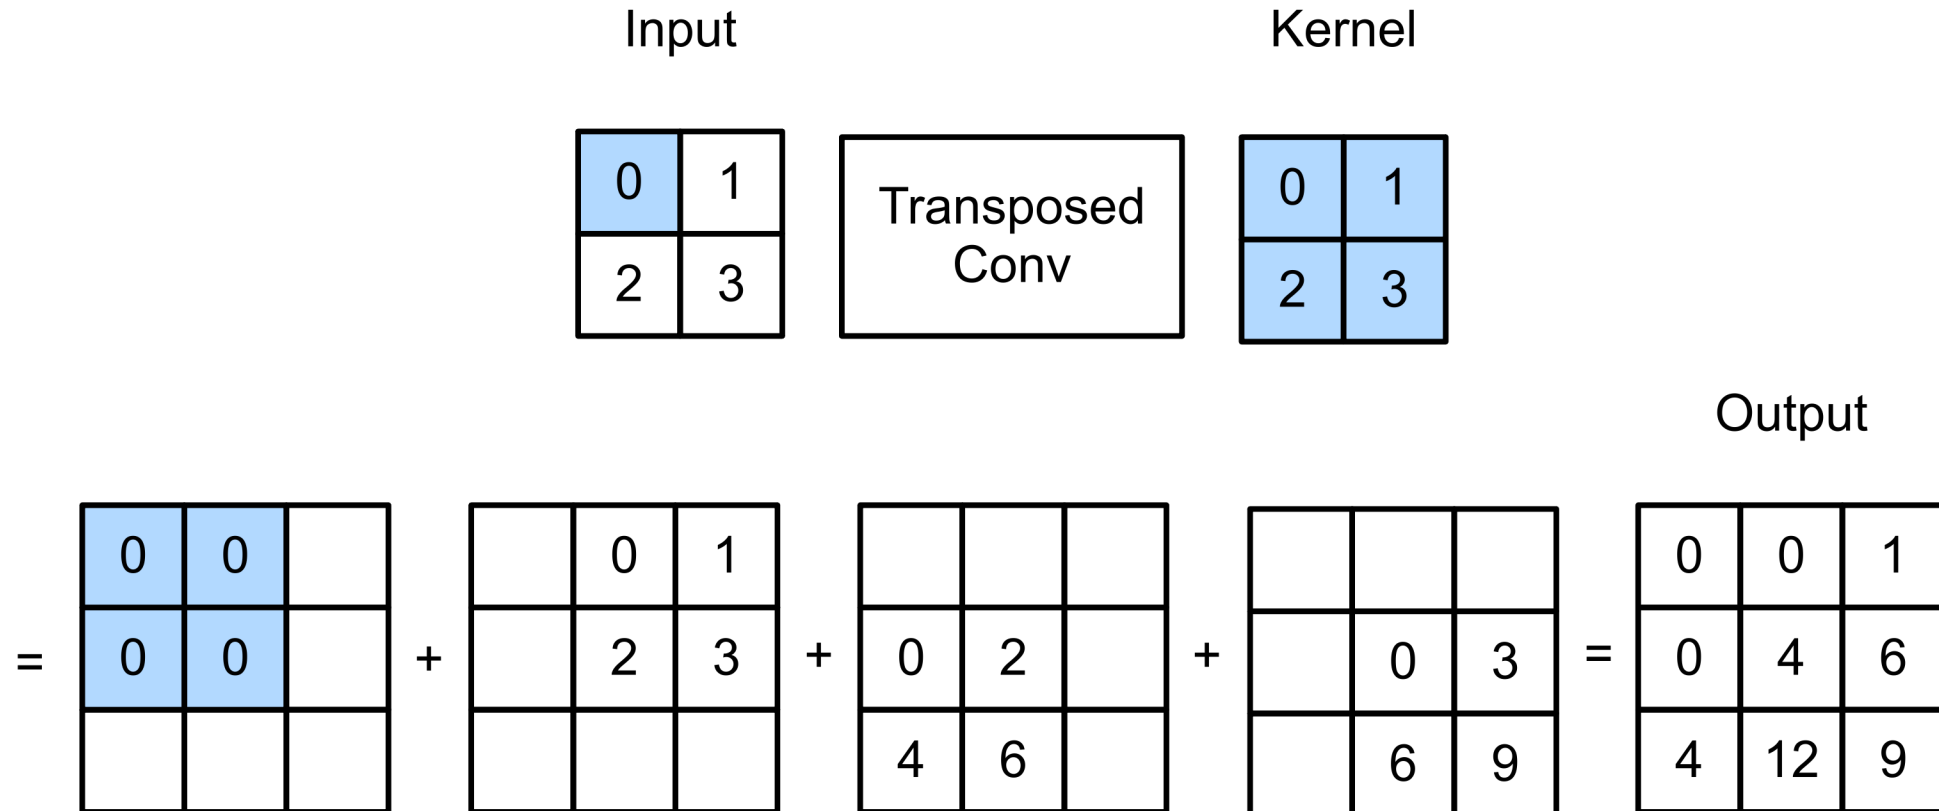
Generator



Discriminator



DCGAN – Transposed convolution



DCGAN – Transposed convolution

